

巨量資料分析與應用 期末報告

基本資料

班級：商資三甲 分組編號：第 8 組

組長：學號 C108193121 姓名 黃昱綺

組員：學號 C108193104 姓名 鄭云瑄

題目

Heart failure clinical records Data Set

心力衰竭臨床記錄數據集

--資料分析目的說明--

說明本分析的目標或目的

本組利用「心力衰竭臨床記錄數據集」，age（年齡）、anaemia（貧血）、high blood pressure（高血壓）、creatinine phosphokinase（肌酐磷酸激酶 CPK）、dabetes（糖尿病）、ejection fraction（射血分數）、platelets（血小板）、sex（性別）、serum creatinine（血清肌酐）、serum sodium（血清鈉）、smoking（吸煙）、time（時間），十二個特徵值預測心力衰竭患者的生存率。

資料集描述

資料集個數：1（以下表格不夠自行拷貝）

總筆數：300

資料集 1 名稱：heart

資料集 1 筆數：300

檔名 1：heart_failure_clinical_records_dataset.csv

來源：

https://archive.ics.uci.edu/ml/machine-learning-databases/00519/heart_failure_clinical_records_dataset.csv

欄位名稱	欄位描述（包括編碼格式）
age 年齡	患者年齡（歲）。
anaemia 貧血	紅細胞或血紅蛋白減少（布林值）。
high blood pressure 高血壓	如果患者患有高血壓（布林值）。
creatinine phosphokinase (CPK) 肌酐磷酸激酶	血液中 CPK 酶的水平（mcg/L）。
dabetes 糖尿病	如果患者患有糖尿病（布林值）。
ejection fraction 射血分數	每次收縮時離開心臟的血液百分比（百分比）。
patelets 血小板	血液中的血小板（千血小板/mL）。
sex 性別	女性或男性（二元）。
serum creatinine 血清肌酐	血液中的血清肌酐水平（mg/dL）。
serum sodium 血清鈉	血液中的血清鈉水平（mEq/L）。
smoking 吸菸	患者是否吸煙（布林值）。
time 時間	觀察期（天）。
[target] death event [目標] 死亡事件	如果患者在觀察期內死亡（布林值）。

分析策略說明

說明通過哪些步驟或流程？使用哪些資料集、哪些欄位、用何種方法來分析？

使用方法：迴歸分析。

1. 先進入 Ubuntu 作業系統並啟動 hadoop。
2. 下載心力衰竭臨床記錄數據集並上傳至 hdfs。
3. 進入 spark 環境並匯入程式庫。
4. 讀取文本資料，將以逗號分割資料成陣列型態且首行（欄位名稱）進行刪除，並轉為浮點數。
5. 建立欄位對照表，有以下欄位：age（年齡）、anaemia（貧血）、high blood pressure（高血壓）、creatinine phosphokinase（肌酐磷酸激酶 CPK）、dabetes（糖尿病）、ejection fraction（射血分數）、platelets（血小板）、sex（性別）、serum creatinine（血清肌酐）、serum sodium（血清鈉）、smoking（吸煙）、time（時間）。
6. 列印出數據集的基本統計值。
7. 建立資料類別後進行資料標準化。
8. 切割資料集，將資料用 7:2:1 的比例，分成訓練集、驗證集與測試集。
9. 以步進 0.1 和迭代次數 1000 開始訓練模型。
10. 列印出各特徵權重。
11. 求驗證集預測值。
12. 寫 getMSE 副程式取得真實評分與預測評分，合併後算出 MSE。
13. 寫 modelAndMSE 副程式訓練各參數組合（步進+迭代次數）的模型求 MSE 值並選取最低 MSE 值的模型。
14. 列出各特徵變數權重。

程式碼

- 請使用 Consolas 字型

進入 Ubuntu 環境	
start-all.sh	啟動 hadoop，若已經啟動，則不需要再執行此指令。
cd ~	回到家目錄。
wget https://archive.ics.uci.edu/ml/machine-learning-databases/00519/heart_failure_clinical_records_dataset.csv	下載心力衰竭臨床記錄數據集。
hdfs dfs -mkdir /heart	在 hdfs 上建立目錄。
hdfs dfs -put ~/heart_failure_clinical_records_dataset.csv /heart	將檔案上傳 hdfs。
hdfs dfs -ls /heart	查看 hdfs 資料。
進入 spark-shell 環境，開始處理	
spark-shell	進入 spark-shell
import org.apache.spark.mllib.linalg.Vectors import org.apache.spark.mllib.regression.LabeledPoint import org.apache.spark.mllib.regression.LinearRegressionWithSGD import org.apache.spark.mllib.feature.StandardScaler	匯入四個必要的程式庫。
val rawData=sc.textFile("/home/mis/heart_failure_clinical_records_dataset.csv") val heartStringRDD=rawData.map(line => line.split(",")).mapPartitionsWithIndex { (index,lines) => if (index==0) lines.drop(1) else lines } val	讀入文本資料，去除首行（欄位名稱），並轉成浮點數。

<pre>heartRDD=heartStringRDD.map(x=>x.map(x=>x .toDouble)) heartRDD.count heartRDD.first</pre>	
<pre>val colNameMap=Map(0->"age", 1->"anaemia", 2->"creatinine_phosphokinase", 3->"diabetes", 4->"ejection_fraction", 5->"high_blood_pressure", 6->"platelets", 7->"serum_creatinine", 8->"serum_sodium", 9->"sex", 10->"smoking", 11->"time", 12->"DEATH_EVENT")</pre>	建立欄位對照表。
基本統計資料	
<pre>def printStats(rdd: org.apache.spark.rdd.RDD[Array[Double]]) { for(i<-0 until rdd.first.length) { val name=colNameMap(i) val s=rdd.map(x=>x(i)).stats val sp="".padTo(25-name.length," ").mkString("") println(f"\$name\$sp\$s") } } printStats(heartRDD)</pre>	列印資料集的基本統計值。
LabeledPoint	
<pre>val heartLP=heartRDD.map(line=>new LabeledPoint(line.last, Vectors.dense(line.init))) heartLP.first</pre>	建立 LabeledPoint 資料類別。

<pre>heartLP.first.label heartLP.first.features</pre>	
資料標準化 (standardize)	
<pre>val heartScaler = new StandardScaler(withMean = true, withStd = true).fit(heartLP.map(x=>x.features))</pre>	建立 heartScaler 物件, 配適 heartLP 的尺度。
<pre>val scaledHeartLP=heartLP.map(x => new LabeledPoint(x.label, heartScaler.transform(x.features))) printStats(scaledHeartLP.map(x=>x.feature s.toArray))</pre>	進行資料標準化。
切割資料集	
<pre>val heartSplit = scaledHeartLP.randomSplit(Array(0.7, 0.2, 0.1),1688) val heartTrainSet = heartSplit(0) val heartValidSet = heartSplit(1) val heartTestSet = heartSplit(2) heartTrainSet.count heartValidSet.count heartTestSet.count heartTrainSet.cache</pre>	切割資料集。
模型訓練	
<pre>val stepSize=0.1 val numOfIter = 1000 val heartModelInst = new LinearRegressionWithSGD().setIntercept(tr ue) heartModelInst.optimizer.setNumIterations (numOfIter).setStepSize(stepSize) val heartModel=heartModelInst.run(heartTrainS et)</pre>	以步進 0.1 和迭代次數 1000 訓練模型。

heartModel.weights	
<pre>def printWeights(w: org.apache.spark.mllib.linalg.Vector) { val pw=colNameMap.toArray.sortBy(x=>x._1).map (x=>x._2).zip(w.toArray).sortBy(x=>x._2). reverse pw.foreach{ case (name, wgt) => val sp="".padTo(25-name.length," ").mkString("") println(f"\$name\$sp\$wgt") } } printWeights(heartModel.weights)</pre>	列印各特徵權重的副程式。
預測與模式評估	
<pre>val heartPred=heartModel.predict(heartValidSe t.map(x=>x.features)) heartPred.first heartValidSet.first</pre>	求驗證集預測值。
<pre>def getMSE(model: org.apache.spark.mllib.regression.LinearR egressionModel, dataset: org.apache.spark.rdd.RDD[org.apache.spark .mllib.regression.LabeledPoint])={ val heartReal=dataset.map(x=>x.label) // 取得真實評分 val heartPred=model.predict(dataset.map(x=>x. features)) // 預測評分 val realWithPred=heartReal.zip(heartPred) // 真實與預測分數合併 val MSE=realWithPred.map{case (real, pred)=>math.pow(real-pred,2)}.mean() MSE // 最後結果 MSE } val MSE=getMSE(heartModel, heartValidSet)</pre>	計算 MSE 的副程式，輸入一個已經訓練好的模型和資料集（通常是驗證資料集）。

println("MSE: "+ MSE)	
最佳化參數的最佳組合	
<pre>def modelAndMSE(stepSize:Double, numOfIter:Int, dataset: org.apache.spark.rdd.RDD[org.apache.spark .mllib.regression.LabeledPoint]) { // 訓練模型 val heartModelInst = new LinearRegressionWithSGD().setIntercept(tr ue) heartModelInst.optimizer.setNumIterations (numOfIter).setStepSize(stepSize) val heartModel=heartModelInst.run(heartTrainS et) // 求預測值 val heartPred=heartModel.predict(dataset.map(x=>x.features)) val heartReal=dataset.map(x=>x.label) // 計算 MSE val MSE=getMSE(heartModel, dataset) println("MSE: "+ MSE+"<=="+"stepSize+"+"/"+numOfIter) }</pre>	<p>訓練某一參數組合（步進+迭代次數）的模型，並計算 MSE。</p>
<pre>for(i<-Array(0.1,0.2,0.3,0.4,0.5,0.6); j<-Array(100,150,200,250,300)) { modelAndMSE(i,j, heartValidSet) }</pre>	<p>求取各超參數組合的 MSE。</p>
<pre>val stepSize=0.5 val numOfIter=100 val heartModelInst = new LinearRegressionWithSGD().setIntercept(tr ue) heartModelInst.optimizer.setNumIterations (numOfIter).setStepSize(stepSize)</pre>	<p>使用最佳超參數組合訓練模型。</p>

<pre>val heartModel=heartModelInst.run(heartTrainS et)</pre>	
<pre>val MSE=getMSE(heartModel , heartTestSet) println("MSE: "+ MSE)</pre>	求最佳模型的 MSE。
<pre>printWeights(heartModel.weights)</pre>	列印各特徵變數權重。

HardCopy 執行畫面

```
mis@master: ~  
mis@master:~$ start-all.sh  
This script is Deprecated. Instead use start-dfs.sh and start-yarn.sh  
Starting namenodes on [master]  
master: starting namenode, logging to /usr/local/hadoop/logs/hadoop-mis-namenode-master.out  
slave1: starting datanode, logging to /usr/local/hadoop/logs/hadoop-mis-datanode-slave1.out  
Starting secondary namenodes [0.0.0.0]  
0.0.0.0: starting secondarynamenode, logging to /usr/local/hadoop/logs/hadoop-mis-secondarynamenode-master.out  
starting yarn daemons  
starting resourcemanager, logging to /usr/local/hadoop/logs/yarn-mis-resourcemanager-master.out  
slave1: starting nodemanager, logging to /usr/local/hadoop/logs/yarn-mis-nodemanager-slave1.out  
mis@master:~$ cd ~  
mis@master:~$ wget https://archive.ics.uci.edu/ml/machine-learning-databases/00519/heart_failure_clinical_records_dataset.csv  
--2022-01-13 13:58:45-- https://archive.ics.uci.edu/ml/machine-learning-databases/00519/heart_failure_clinical_records_dataset.csv  
正在查找主機 archive.ics.uci.edu (archive.ics.uci.edu)... 128.195.10.252  
正在連接 archive.ics.uci.edu (archive.ics.uci.edu)[128.195.10.252]:443... 連上了。  
已送出 HTTP 要求，正在等候回應... 200 OK  
長度: 12239 (12K) [application/x-httpd-php]  
Saving to: 'heart_failure_clinical_records_dataset.csv.1'  
  
heart_failure_clinical_r 100%[=====] 11.95K --.-KB/s in 0.006s  
  
2022-01-13 13:58:51 (1.96 MB/s) - 'heart_failure_clinical_records_dataset.csv.1' saved [12239/12239]  
  
mis@master:~$ hdfs dfs -mkdir /heart  
mis@master:~$ hdfs dfs -put ~/heart_failure_clinical_records_dataset.csv /heart  
mis@master:~$ hdfs dfs -ls /heart  
Found 1 items  
-rw-r--r-- 1 mis supergroup 12239 2022-01-13 13:59 /heart/heart_failure_clinical_records_dataset.csv
```

```
mis@master: ~  
mis@master:~$ spark-shell  
22/01/13 14:05:25 WARN NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable  
Welcome to  
  
      _ _ _ _ _  
     / _ _ _ _ \  
    / _ _ _ _ \  
   / _ _ _ _ \  
  / _ _ _ _ \  
 / _ _ _ _ \  
/_ _ _ _ _ \  
version 1.6.2  
  
Using Scala version 2.10.5 (OpenJDK 64-Bit Server VM, Java 1.8.0_292)  
Type in expressions to have them evaluated.  
Type :help for more information.  
Spark context available as sc.  
22/01/13 14:05:39 WARN Connection: BoneCP specified but not present in CLASSPATH (or one of dependencies)  
22/01/13 14:05:40 WARN Connection: BoneCP specified but not present in CLASSPATH (or one of dependencies)  
22/01/13 14:05:50 WARN ObjectStore: Version information not found in metastore. hive.metastore.schema.verification is not enabled so recording the schema version 1.2.0  
22/01/13 14:05:50 WARN ObjectStore: Failed to get database default, returning NoSuchObjectException  
22/01/13 14:05:54 WARN Connection: BoneCP specified but not present in CLASSPATH (or one of dependencies)  
22/01/13 14:05:55 WARN Connection: BoneCP specified but not present in CLASSPATH (or one of dependencies)  
SQL context available as sqlContext.  
  
scala> import org.apache.spark.mllib.linalg.Vectors  
import org.apache.spark.mllib.linalg.Vectors  
  
scala> import org.apache.spark.mllib.regression.LabeledPoint  
import org.apache.spark.mllib.regression.LabeledPoint  
  
scala> import org.apache.spark.mllib.regression.LinearRegressionWithSGD  
import org.apache.spark.mllib.regression.LinearRegressionWithSGD
```

```

mis@master: ~
scala> import org.apache.spark.mllib.linalg.Vectors
import org.apache.spark.mllib.linalg.Vectors

scala> import org.apache.spark.mllib.regression.LabeledPoint
import org.apache.spark.mllib.regression.LabeledPoint

scala> import org.apache.spark.mllib.regression.LinearRegressionWithSGD
import org.apache.spark.mllib.regression.LinearRegressionWithSGD

scala> import org.apache.spark.mllib.feature.StandardScaler
import org.apache.spark.mllib.feature.StandardScaler

scala> val rawData=sc.textFile("/home/mis/heart_failure_clinical_records_dataset.csv")
rawData: org.apache.spark.rdd.RDD[String] = /home/mis/heart_failure_clinical_records_dataset.csv MapPartit
ionsRDD[1] at textFile at <console>:31

scala>

scala> val heartStringRDD=rawData.map(line => line.split(",")).mapPartitionsWithIndex { (index,lines) => i
f (index==0) lines.drop(1) else lines }
heartStringRDD: org.apache.spark.rdd.RDD[Array[String]] = MapPartitionsRDD[3] at mapPartitionsWithIndex at
<console>:33

scala>

scala> val heartRDD=heartStringRDD.map(x=>x.map(x=>x.toDouble))
heartRDD: org.apache.spark.rdd.RDD[Array[Double]] = MapPartitionsRDD[4] at map at <console>:35

scala>

scala> heartRDD.count

mis@master: ~
scala> heartRDD.count
res0: Long = 299

scala> heartRDD.first
res1: Array[Double] = Array(75.0, 0.0, 582.0, 0.0, 20.0, 1.0, 265000.0, 1.9, 130.0, 1.0, 0.0, 4.0, 1.0)

scala> val colNameMap=Map(
  | 0->"age",
  | 1->"anaemia",
  | 2->"creatinine_phosphokinase",
  | 3->"diabetes",
  | 4->"ejection_fraction",
  | 5->"high_blood_pressure",
  | 6->"platelets",
  | 7->"serum_creatinine",
  | 8->"serum_sodium",
  | 9->"sex",
  | 10->"smoking",
  | 11->"time",
  | 12->"DEATH_EVENT"
  | )
colNameMap: scala.collection.immutable.Map[Int,String] = Map(0 -> age, 5 -> high_blood_pressure, 10 -> smo
king, 1 -> anaemia, 6 -> platelets, 9 -> sex, 2 -> creatinine_phosphokinase, 12 -> DEATH_EVENT, 7 -> serum
_creatinine, 3 -> diabetes, 11 -> time, 8 -> serum_sodium, 4 -> ejection_fraction)

scala> def printStats(rdd: org.apache.spark.rdd.RDD[Array[Double]]) {
  |   for(i<-0 until rdd.first.length) {
  |     val name=colNameMap(i)
  |     val s=rdd.map(x=>x(i)).stats
  |     val sp="".padTo(25-name.length," ").mkString("")
  |     println(f"$name$sp$")
  |   }
}

```

```

mis@master: ~
scala> def printStats(rdd: org.apache.spark.rdd.RDD[Array[Double]]) {
  |   for(i<-0 until rdd.first.length) {
  |     val name=colNameMap(i)
  |     val s=rdd.map(x=>x(i)).stats
  |     val sp="".padTo(25-name.length," ").mkString("")
  |     println(f"$name$sp$s")
  |   }
  | }
printStats: (rdd: org.apache.spark.rdd.RDD[Array[Double]])Unit

scala>

scala> printStats(heartRDD)
age (count: 299, mean: 60.833893, stdev: 11.874901, max: 95.000000, min: 40.000000)
anaemia (count: 299, mean: 0.431438, stdev: 0.495277, max: 1.000000, min: 0.000000)
creatinine_phosphokinase (count: 299, mean: 581.839465, stdev: 968.663967, max: 7861.000000, min: 23.000000)
diabetes (count: 299, mean: 0.418060, stdev: 0.493240, max: 1.000000, min: 0.000000)
ejection_fraction (count: 299, mean: 38.083612, stdev: 11.815033, max: 80.000000, min: 14.000000)
high_blood_pressure (count: 299, mean: 0.351171, stdev: 0.477336, max: 1.000000, min: 0.000000)
platelets (count: 299, mean: 263358.029264, stdev: 97640.547655, max: 850000.000000, min: 25100.000000)
serum_creatinine (count: 299, mean: 1.393880, stdev: 1.032779, max: 9.400000, min: 0.500000)
serum_sodium (count: 299, mean: 136.625418, stdev: 4.405092, max: 148.000000, min: 113.000000)
sex (count: 299, mean: 0.648829, stdev: 0.477336, max: 1.000000, min: 0.000000)
smoking (count: 299, mean: 0.321070, stdev: 0.466888, max: 1.000000, min: 0.000000)
time (count: 299, mean: 130.260870, stdev: 77.484310, max: 285.000000, min: 4.000000)
DEATH_EVENT (count: 299, mean: 0.321070, stdev: 0.466888, max: 1.000000, min: 0.000000)

scala> val heartLP=heartRDD.map(line=>new LabeledPoint(line.last, Vectors.dense(line.init)))
heartLP: org.apache.spark.rdd.RDD[org.apache.spark.mllib.regression.LabeledPoint] = MapPartitionsRDD[31] a

mis@master: ~

scala> val heartLP=heartRDD.map(line=>new LabeledPoint(line.last, Vectors.dense(line.init)))
heartLP: org.apache.spark.rdd.RDD[org.apache.spark.mllib.regression.LabeledPoint] = MapPartitionsRDD[31] a
t map at <console>:37

scala> heartLP.first
res3: org.apache.spark.mllib.regression.LabeledPoint = (1.0,[75.0,0.0,582.0,0.0,20.0,1.0,265000.0,1.9,130.0,1.0,0.0,4.0])

scala> heartLP.first.label
res4: Double = 1.0

scala> heartLP.first.features
res5: org.apache.spark.mllib.linalg.Vector = [75.0,0.0,582.0,0.0,20.0,1.0,265000.0,1.9,130.0,1.0,0.0,4.0]

scala>

scala> val heartScaler = new StandardScaler(withMean = true, withStd = true).fit(heartLP.map(x=>x.features))
heartScaler: org.apache.spark.mllib.feature.StandardScalerModel = org.apache.spark.mllib.feature.StandardScalerModel@321d3e08

scala> val scaledHeartLP=heartLP.map(x => new LabeledPoint(x.label, heartScaler.transform(x.features)))
scaledHeartLP: org.apache.spark.rdd.RDD[org.apache.spark.mllib.regression.LabeledPoint] = MapPartitionsRDD[34] at map at <console>:41

scala> printStats(scaledHeartLP.map(x=>x.features.toArray))
age (count: 299, mean: -0.000000, stdev: 0.998326, max: 2.872354, min: -1.751511)
anaemia (count: 299, mean: 0.000000, stdev: 0.998326, max: 1.146046, min: -0.869647)
creatinine_phosphokinase (count: 299, mean: 0.000000, stdev: 0.998326, max: 7.502063, min: -0.575952)
diabetes (count: 299, mean: 0.000000, stdev: 0.998326, max: 1.177856, min: -0.846161)
ejection_fraction (count: 299, mean: -0.000000, stdev: 0.998326, max: 3.541779, min: -2.034976)

```

```

mis@master: ~
scala> printStats(scaledHeartLP.map(x=>x.features.toArray))
age (count: 299, mean: -0.000000, stdev: 0.998326, max: 2.872354, min: -1.751511)
anaemia (count: 299, mean: 0.000000, stdev: 0.998326, max: 1.146046, min: -0.869647)
creatinine_phosphokinase (count: 299, mean: 0.000000, stdev: 0.998326, max: 7.502063, min: -0.575952)
diabetes (count: 299, mean: 0.000000, stdev: 0.998326, max: 1.177856, min: -0.846161)
ejection_fraction (count: 299, mean: -0.000000, stdev: 0.998326, max: 3.541779, min: -2.034976)
high_blood_pressure (count: 299, mean: -0.000000, stdev: 0.998326, max: 1.356997, min: -0.734457)
platelets (count: 299, mean: 0.000000, stdev: 0.998326, max: 5.998124, min: -2.436071)
serum_creatinine (count: 299, mean: -0.000000, stdev: 0.998326, max: 7.739045, min: -0.864061)
serum_sodium (count: 299, mean: -0.000000, stdev: 0.998326, max: 2.577822, min: -5.354230)
sex (count: 299, mean: -0.000000, stdev: 0.998326, max: 0.734457, min: -1.356997)
smoking (count: 299, mean: 0.000000, stdev: 0.998326, max: 1.451727, min: -0.686531)
time (count: 299, mean: 0.000000, stdev: 0.998326, max: 1.993696, min: -1.626775)

scala> val heartSplit = scaledHeartLP.randomSplit(Array(0.7, 0.2, 0.1),1688)
heartSplit: Array[org.apache.spark.rdd.RDD[org.apache.spark.mllib.regression.LabeledPoint]] = Array(MapPartitionsRDD[60] at randomSplit at <console>:43, MapPartitionsRDD[61] at randomSplit at <console>:43, MapPartitionsRDD[62] at randomSplit at <console>:43)

scala> val heartTrainSet = heartSplit(0)
heartTrainSet: org.apache.spark.rdd.RDD[org.apache.spark.mllib.regression.LabeledPoint] = MapPartitionsRDD[60] at randomSplit at <console>:43

scala> val heartValidSet = heartSplit(1)
heartValidSet: org.apache.spark.rdd.RDD[org.apache.spark.mllib.regression.LabeledPoint] = MapPartitionsRDD[61] at randomSplit at <console>:43

scala> val heartTestSet = heartSplit(2)
heartTestSet: org.apache.spark.rdd.RDD[org.apache.spark.mllib.regression.LabeledPoint] = MapPartitionsRDD[62] at randomSplit at <console>:43

mis@master: ~
scala> val heartTestSet = heartSplit(2)
heartTestSet: org.apache.spark.rdd.RDD[org.apache.spark.mllib.regression.LabeledPoint] = MapPartitionsRDD[62] at randomSplit at <console>:43

scala> heartTrainSet.count
res7: Long = 202

scala> heartValidSet.count
res8: Long = 54

scala> heartTestSet.count
res9: Long = 43

scala> heartTrainSet.cache
res10: heartTrainSet.type = MapPartitionsRDD[60] at randomSplit at <console>:43

scala> val stepSize=0.1
stepSize: Double = 0.1

scala> val numOfIter = 1000
numOfIter: Int = 1000

scala> val heartModelInst = new LinearRegressionWithSGD().setIntercept(true)
heartModelInst: org.apache.spark.mllib.regression.LinearRegressionWithSGD = org.apache.spark.mllib.regression.LinearRegressionWithSGD@4b7cb8a0

scala> heartModelInst.optimizer.setNumIterations(numOfIter).setStepSize(stepSize)
res11: heartModelInst.optimizer.type = org.apache.spark.mllib.optimization.GradientDescent@26327339

scala> val heartModel=heartModelInst.run(heartTrainSet)
22/01/13 14:18:34 WARN BLAS: Failed to load implementation from: com.github.fommil.netlib.NativeSystemBLAS

```



```

mis@master: ~
scala> val heartModel=heartModelInst.run(heartTrainSet)
22/01/13 14:18:34 WARN BLAS: Failed to load implementation from: com.github.fommil.netlib.NativeSystemBLAS
22/01/13 14:18:34 WARN BLAS: Failed to load implementation from: com.github.fommil.netlib.NativeRefBLAS
heartModel: org.apache.spark.mllib.regression.LinearRegressionModel = org.apache.spark.mllib.regression.LinearRegressionModel: intercept = 0.4238865944504079, numFeatures = 12

scala> heartModel.weights
res12: org.apache.spark.mllib.linalg.Vector = [0.062322493078632475,-0.010363033957126044,0.026406997937481715,0.00884144057118402,-0.08914726339228686,0.0121666389431587,-0.005024070891089872,0.07763039181358414,-0.030071342377087043,-0.03098151872337847,-0.002904035263979059,-0.18083004200203062]

scala> def printWeights(w: org.apache.spark.mllib.linalg.Vector) {
    | val pw=colNameMap.toArray.sortBy(x=>x._1).map(x=>x._2).zip(w.toArray).sortBy(x=>x._2).reverse
    | pw.foreach{ case (name, wgt) =>
    |   | val sp="".padTo(25-name.length," ").mkString("")
    |   | println(f"$name$sp$wgt")
    | }
    | }
printWeights: (w: org.apache.spark.mllib.linalg.Vector)Unit

scala> printWeights(heartModel.weights)
serum_creatinine      0.07763039181358414
age                   0.062322493078632475
creatinine_phosphokinase 0.026406997937481715
high_blood_pressure   0.0121666389431587
diabetes              0.00884144057118402
smoking               -0.002904035263979059

```

```

mis@master: ~
scala> printWeights(heartModel.weights)
serum_creatinine      0.07763039181358414
age                   0.062322493078632475
creatinine_phosphokinase 0.026406997937481715
high_blood_pressure   0.0121666389431587
diabetes              0.00884144057118402
smoking               -0.002904035263979059
platelets             -0.005024070891089872
anaemia               -0.010363033957126044
serum_sodium          -0.030071342377087043
sex                   -0.03098151872337847
ejection_fraction     -0.08914726339228686
time                  -0.18083004200203062

scala> val heartPred=heartModel.predict(heartValidSet.map(x=>x.features))
heartPred: org.apache.spark.rdd.RDD[Double] = MapPartitionsRDD[290] at mapPartitions at GeneralizedLinearAlgorithm.scala:69

scala> heartPred.first
res14: Double = 0.8403665752548327

scala> heartValidSet.first
res15: org.apache.spark.mllib.regression.LabeledPoint = (1.0,[2.031651527399793,-0.8696468582568094,-0.5759522261296451,-0.8461608356372154,0.5844090437045386,-0.7344569121685459,0.9881164030309463,1.5525420748331058,-1.0482587813127016,0.7344569121685459,-0.6865309732064924,-1.3175534772023767])

scala> def getMSE(model: org.apache.spark.mllib.regression.LinearRegressionModel, dataset: org.apache.spar

```

```

mis@master: ~
scala> def getMSE(model: org.apache.spark.mllib.regression.LinearRegressionModel, dataset: org.apache.spark.rdd.RDD[org.apache.spark.mllib.regression.LabeledPoint])={
|   val heartReal=dataset.map(x=>x.label) // 取得真實評分
|   val heartPred=model.predict(dataset.map(x=>x.features)) // 預測評分
|   val realWithPred=heartReal.zip(heartPred) // 真實與預測分數合併
|   val MSE=realWithPred.map{case (real, pred)=>math.pow(real-pred,2)}.mean()
|   MSE // 最後結果MSE
| }
getMSE: (model: org.apache.spark.mllib.regression.LinearRegressionModel, dataset: org.apache.spark.rdd.RDD[org.apache.spark.mllib.regression.LabeledPoint])Double

scala> val MSE=getMSE(heartModel, heartValidSet)
MSE: Double = 0.11923800048370291

scala> println("MSE: " + MSE)
MSE: 0.11923800048370291

scala> def modelAndMSE(stepSize:Double, numOfIter:Int, dataset: org.apache.spark.rdd.RDD[org.apache.spark.mllib.regression.LabeledPoint]) {
|   // 訓練模型
|   val heartModelInst = new LinearRegressionWithSGD().setIntercept(true)
|   heartModelInst.optimizer.setNumIterations(numOfIter).setStepSize(stepSize)
|   val heartModel=heartModelInst.run(heartTrainSet)
|
|   // 求預測值
|   val heartPred=heartModel.predict(dataset.map(x=>x.features))
|   val heartReal=dataset.map(x=>x.label)
| }

```

```

mis@master: ~
scala> def modelAndMSE(stepSize:Double, numOfIter:Int, dataset: org.apache.spark.rdd.RDD[org.apache.spark.mllib.regression.LabeledPoint]) {
|   // 訓練模型
|   val heartModelInst = new LinearRegressionWithSGD().setIntercept(true)
|   heartModelInst.optimizer.setNumIterations(numOfIter).setStepSize(stepSize)
|   val heartModel=heartModelInst.run(heartTrainSet)
|
|   // 求預測值
|   val heartPred=heartModel.predict(dataset.map(x=>x.features))
|   val heartReal=dataset.map(x=>x.label)
|
|   // 計算MSE
|   val MSE=getMSE(heartModel, dataset)
|   println("MSE: " + MSE+"<=="+stepSize+"/"+numOfIter)
| }
modelAndMSE: (stepSize: Double, numOfIter: Int, dataset: org.apache.spark.rdd.RDD[org.apache.spark.mllib.regression.LabeledPoint])Unit

scala> for(i<-Array(0.1,0.2,0.3,0.4,0.5,0.6); j<-Array(100,150,200,250,300)) {
|   modelAndMSE(i,j, heartValidSet)
| }
MSE: 0.1231337560078758<==0.1/100
MSE: 0.11923800048370291<==0.1/150
MSE: 0.11923800048370291<==0.1/200
MSE: 0.11923800048370291<==0.1/250
MSE: 0.11923800048370291<==0.1/300

```

```

mis@master: ~
scala> for(i<-Array(0.1,0.2,0.3,0.4,0.5,0.6); j<-Array(100,150,200,250,300)) {
|   modelAndMSE(i,j, heartValidSet)
| }
MSE: 0.1231337560078758<==0.1/100
MSE: 0.11923800048370291<==0.1/150
MSE: 0.11923800048370291<==0.1/200
MSE: 0.11923800048370291<==0.1/250
MSE: 0.11923800048370291<==0.1/300
MSE: 0.10480198450793221<==0.2/100
MSE: 0.10480198450793221<==0.2/150
MSE: 0.10480198450793221<==0.2/200
MSE: 0.10480198450793221<==0.2/250
MSE: 0.10480198450793221<==0.2/300
MSE: 0.10308027481763134<==0.3/100
MSE: 0.10308027481763134<==0.3/150
MSE: 0.10308027481763134<==0.3/200
MSE: 0.10308027481763134<==0.3/250
MSE: 0.10308027481763134<==0.3/300
MSE: 0.10273310302866487<==0.4/100
MSE: 0.10273310302866487<==0.4/150
MSE: 0.10273310302866487<==0.4/200
MSE: 0.10273310302866487<==0.4/250
MSE: 0.10273310302866487<==0.4/300
MSE: 0.10270183551648068<==0.5/100
MSE: 0.10270183551648068<==0.5/150
MSE: 0.10270183551648068<==0.5/200
MSE: 0.10270183551648068<==0.5/250

```

```

mis@master: ~
MSE: 0.10480198450793221<==0.2/150
MSE: 0.10480198450793221<==0.2/200
MSE: 0.10480198450793221<==0.2/250
MSE: 0.10480198450793221<==0.2/300
MSE: 0.10308027481763134<==0.3/100
MSE: 0.10308027481763134<==0.3/150
MSE: 0.10308027481763134<==0.3/200
MSE: 0.10308027481763134<==0.3/250
MSE: 0.10308027481763134<==0.3/300
MSE: 0.10273310302866487<==0.4/100
MSE: 0.10273310302866487<==0.4/150
MSE: 0.10273310302866487<==0.4/200
MSE: 0.10273310302866487<==0.4/250
MSE: 0.10273310302866487<==0.4/300
MSE: 0.10270183551648068<==0.5/100
MSE: 0.10270183551648068<==0.5/150
MSE: 0.10270183551648068<==0.5/200
MSE: 0.10270183551648068<==0.5/250
MSE: 0.10270183551648068<==0.5/300
MSE: 0.10275895378911842<==0.6/100
MSE: 0.10275895378911842<==0.6/150
MSE: 0.10275895378911842<==0.6/200
MSE: 0.10275895378911842<==0.6/250
MSE: 0.10275895378911842<==0.6/300

scala> val stepSize=0.5
stepSize: Double = 0.5

```



```

mis@master: ~
scala> val stepSize=0.5
stepSize: Double = 0.5

scala> val numOfIter=100
numOfIter: Int = 100

scala> val heartModelInst = new LinearRegressionWithSGD().setIntercept(true)
heartModelInst: org.apache.spark.mllib.regression.LinearRegressionWithSGD = org.apache.spark.mllib.regression.LinearRegressionWithSGD@393f9d65

scala> heartModelInst.optimizer.setNumIterations(numOfIter).setStepSize(stepSize)
res18: heartModelInst.optimizer.type = org.apache.spark.mllib.optimization.GradientDescent@6806eae7

scala> val heartModel=heartModelInst.run(heartTrainSet)
heartModel: org.apache.spark.mllib.regression.LinearRegressionModel = org.apache.spark.mllib.regression.LinearRegressionModel: intercept = 0.32850932792440857, numFeatures = 12

scala> val MSE=getMSE(heartModel , heartTestSet)
MSE: Double = 0.12946912443154074

scala> println("MSE: " + MSE)
MSE: 0.12946912443154074

scala> printWeights(heartModel.weights)
serum_creatinine      0.08795089849459757
age                   0.07435519397223693
creatinine_phosphokinase 0.02886542946874809
smoking               0.013658016119084557

```

```

mis@master: ~
scala> heartModelInst.optimizer.setNumIterations(numOfIter).setStepSize(stepSize)
res18: heartModelInst.optimizer.type = org.apache.spark.mllib.optimization.GradientDescent@6806eae7

scala> val heartModel=heartModelInst.run(heartTrainSet)
heartModel: org.apache.spark.mllib.regression.LinearRegressionModel = org.apache.spark.mllib.regression.LinearRegressionModel: intercept = 0.32850932792440857, numFeatures = 12

scala> val MSE=getMSE(heartModel , heartTestSet)
MSE: Double = 0.12946912443154074

scala> println("MSE: " + MSE)
MSE: 0.12946912443154074

scala> printWeights(heartModel.weights)
serum_creatinine      0.08795089849459757
age                   0.07435519397223693
creatinine_phosphokinase 0.02886542946874809
smoking               0.013658016119084557
diabetes              0.01012703020077554
high_blood_pressure   0.003426357786016492
platelets             -0.009806129722302328
anaemia               -0.01203119471944738
serum_sodium          -0.02562457352466533
sex                   -0.04934763385045032
ejection_fraction     -0.10883233979106556
time                  -0.2126882420020556

```